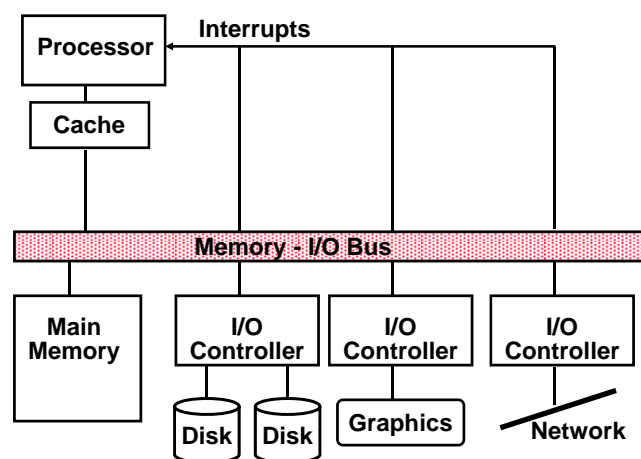


---

## Chapter 4: Input/Output System

### A Typical I/O System



## Input and Output Devices

- ❑ I/O devices are incredibly diverse with respect to
  - Behavior – input, output or storage
  - Partner – human or machine
  - Data rate – the peak rate at which data can be transferred between the I/O device and the main memory or processor

Device	Behavior	Partner	Data rate (Mb/s)
Keyboard	input	human	0.0001
Mouse	input	human	0.0038
Laser printer	output	human	3.2000
Magnetic disk	storage	machine	800.0000-3000.0000
Graphics display	output	human	800.0000-8000.0000
Network/LAN	input or output	machine	100.0000-10000.0000

8 orders of magnitude  
range

## I/O Performance Measures

- ❑ **I/O bandwidth** (throughput) – amount of information that can be input (output) and communicated across an interconnect (e.g., a bus) to the processor/memory (I/O device) per unit time
  1. How much data can we move through the system in a certain time?
  2. How many I/O operations can we do per unit time?
- ❑ **I/O response time** (latency) – the total elapsed time to accomplish an input or output operation
  - An especially important performance metric in real-time systems
- ❑ **Expandability** – is there any easy way to connect another disk to the system?
- ❑ **Resilience** – if this I/O controller (network) fails, is it going to affect the rest of the network?

---

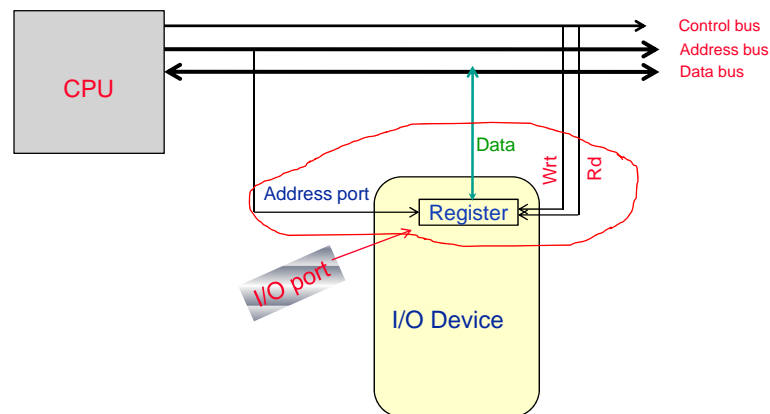
# Input/output Ports

## I/O ports

### □ Communication between CPU and I/O devices

- How does the processor communicate with devices other than main memory?
  - By using Input/output ports
- I/O ports
  - Input port: transfers from external device to CPU
  - Output port: transfers from CPU to external
  - Input/Output ports: transfers in both directions

## I/O ports



## I/O Commands

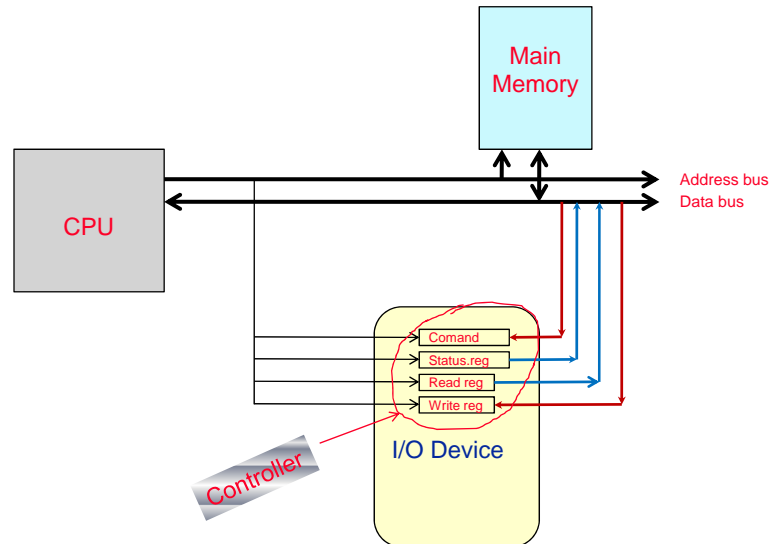
### ❑ I/O devices are managed by I/O controller hardware

- Transfers data to/from device
- Synchronizes operations with software

### ❑ Ports in a I/O controller:

- Command registers
  - Cause device to do something
- Status registers
  - Indicate what the device is doing and occurrence of errors
- Data registers
  - Write: transfer data to a device
  - Read: transfer data from a device

## I/O Commands



## Communication of I/O Devices and Processor

- ❑ User programs (processor in protected mode) are prevented from issuing I/O operations directly because the OS does not provide access to the I/O ports
- ❑ Only when processor is in kernel (supervisor) mode, then the I/O ports can be accessed
- ❑ How the processor directs the I/O devices: through the address space
  1. Memory-mapped I/O
  2. Specific I/O instructions

## Communication of I/O Devices and Processor

### □ How the processor directs the I/O devices

#### 1. Memory-mapped I/O

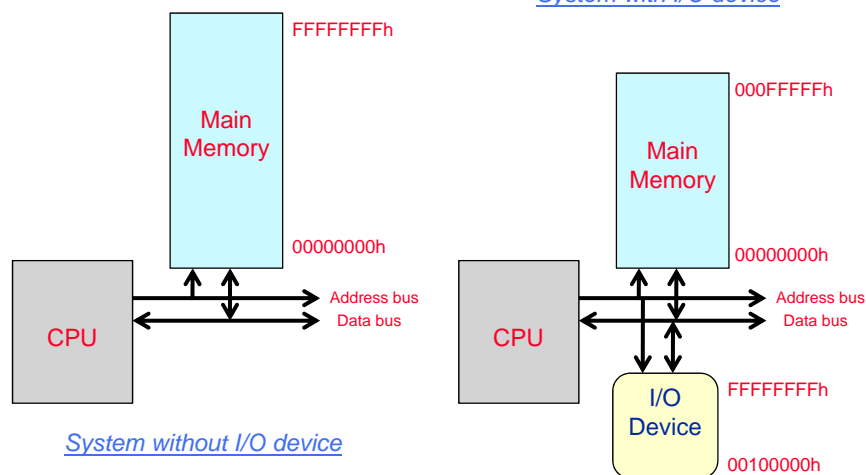
- Portions of the high-order memory address space are assigned to each I/O device
- Read and writes to those memory addresses are interpreted as commands to the I/O devices
- Load/stores to the I/O address space can *only* be done by the OS
- MIPS processor:
  - Load instruction to read from I/O device
    - » i.e. lw \$4, 100(\$5)
  - Store instruction to write to I/O device
    - » i.e. sw \$4, 100(\$5)

## Communication of I/O Devices and Processor

### □ How the processor directs the I/O devices

#### • Memory-mapped I/O (MIPS)

#### System with I/O device



## Communication of I/O Devices and Processor

### □ How the processor directs the I/O devices

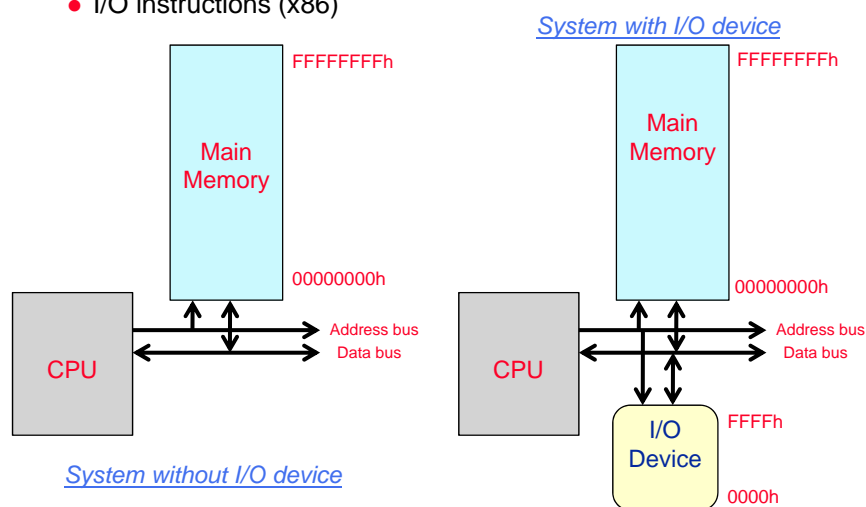
#### 2. I/O instructions

- Separate instructions to access I/O registers
- Can only be executed in kernel mode
- Example: x86:
  - Control signal:  $\overline{IO}/M$  (1 = memory, 0 = IO)
  - Specific Input instruction to read from I/O device
    - » i.e. `in al, 37h`
    - » `P(37) → AL register`
  - Specific Output instruction to write to I/O device
    - » i.e. `out 37h, al`
    - » `AL register → P(37)`

## Communication of I/O Devices and Processor

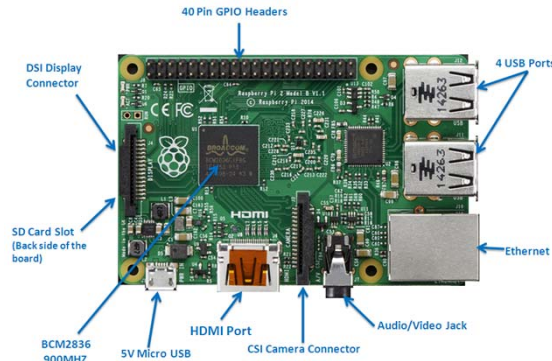
### □ How the processor directs the I/O devices

#### • I/O instructions (x86)



## Raspberry Pi

- ❑ Raspberypi 2 B: a 900MHz quad-core ARM Cortex-A7 CPU 1GB RAM
- ❑ Like the (Pi 1) Model B+, it also has: 100 Base Ethernet, 4 USB ports, 40 GPIO pins, full HDMI port, combined 3.5mm audio jack and composite video, camera interface (CSI), display interface (DSI), micro SD card slot, videoCore IV 3D graphics core.



## Raspberry Pi: GPIO port



Pin No.				
	1	2	3.3V	5V
	3	4	GPIO2	5V
	5	6	GPIO3	GND
	7	8	GPIO4	GPIO14
	9	10	GND	GPIO15
	11	12	GPIO17	GPIO18
	13	14	GPIO27	GND
	15	16	GPIO22	GPIO23
	17	18	3.3V	GPIO24
	19	20	GPIO10	GND
	21	22	GPIO9	GPIO25
	23	24	GPIO11	GPIO8
	25	26	GND	GPIO7

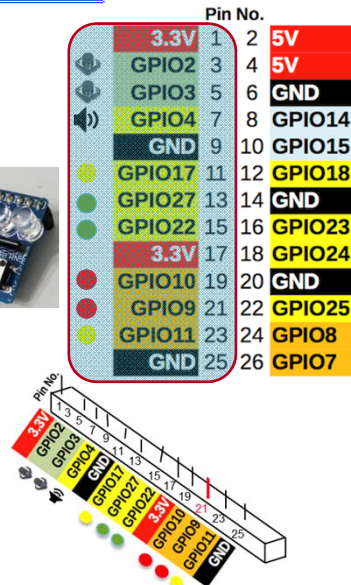
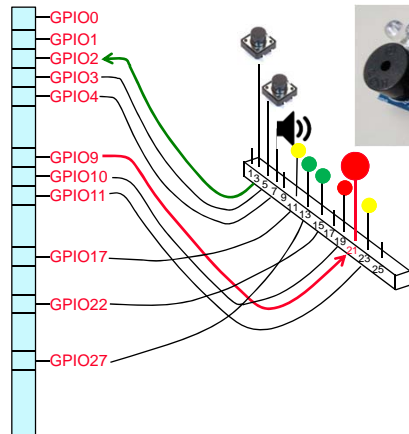
Tx  
Rx

- ❑ Raspberry Pi manages up to 54 pins
- ❑ Only the showed ones are accessible
- ❑ GPIO ports are mapped in memory, starting at 0x3F200000



## GPIO pins

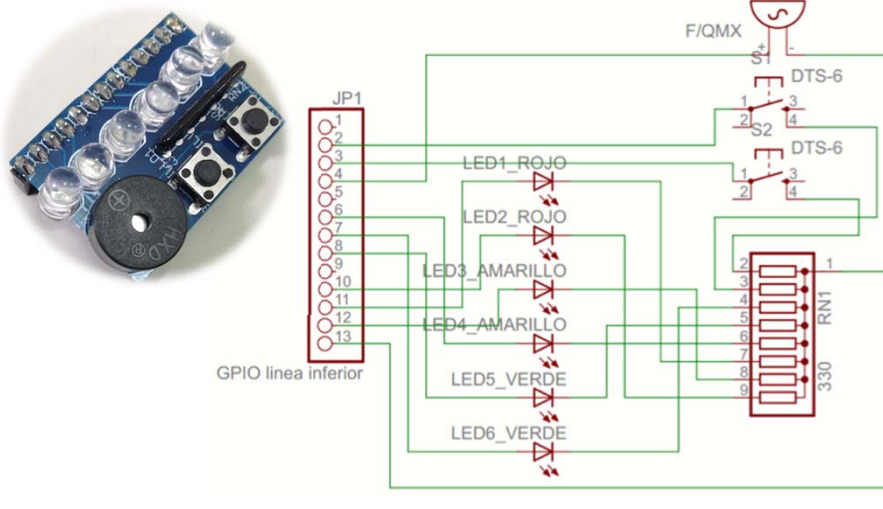
- GPIO: two rows with 13 pines
- We will use the inner row for connecting an external board



Chapter 4.18

Dept. of Comp. Arch., UMA, 2018

## External board connected to GPIO



Chapter 4.19

Dept. of Comp. Arch., UMA, 2018

## GPIO memory mapping

### GPIO ports

- GPFSELn: GPIO Function Select Registers
  - The 54 pins are configured through 6 memory ports, GPSEL0 to GPSEL5
  - Each port defines 10 groups named FSELx, FSEL0 to FSEL9
  - A group consists of 3 bits
  - GPFSEL0 controls GPIO0 to GPIO9, GPFSEL1 controls GPIO10 to GPIO19, ...
- GPSETn: GPIO Pin Output Set Registers
  - GPSET0 sets pins 0 to 31, and GPSET1 sets pins 32 to 53
- GPCLRn: GPIO Pin Output Clear Registers
  - GPCLR0 clears pins 0 to 31, and GPCLR1 clears pins 32 to 53
- A SET or CLR operation in any pin just needs 1 in the corresponding position and only affects that pin (0 means that the pin is not modified)

## GPIO memory mapping

Dir física

3F20 0000	GPFSEL0
3F20 0004	GPFSEL1
3F20 0008	GPFSEL2
3F20 000C	GPFSEL3
3F20 0010	GPFSEL4
3F20 0014	GPFSEL5
3F20 0018	--
3F20 001C	GPSET0
3F20 0020	GPSET1
3F20 0024	--
3F20 0028	GPCLR0
3F20 002C	GPCLR1

Set GPIO9 as Output

X	FSEL9	FSEL8	FSEL7	FSEL6	FSEL5	FSEL4	FSEL3	FSEL2	FSEL1	FSEL0
	001									

Set GPIO9 (turn the led on)

SET31	SET30	SET29	SET28	SET27	SET26	SET25	SET24	SET23	SET22	SET21	SET20	SET19	SET18	SET17	SET16	SET15	SET14	SET13	SET12	SET11	SET10	SET9	SET8	SET7	SET6	SET5	SET4	SET3	SET2	SET1	SET0
																						1									

Clear GPIO9 (turn the led off)

CLR31	CLR30	CLR29	CLR28	CLR27	CLR26	CLR25	CLR24	CLR23	CLR22	CLR21	CLR20	CLR19	CLR18	CLR17	CLR16	CLR15	CLR14	CLR13	CLR12	CLR11	CLR10	CLR9	CLR8	CLR7	CLR6	CLR5	CLR4	CLR3	CLR2	CLR1	CLR0
																						1									

32 bits

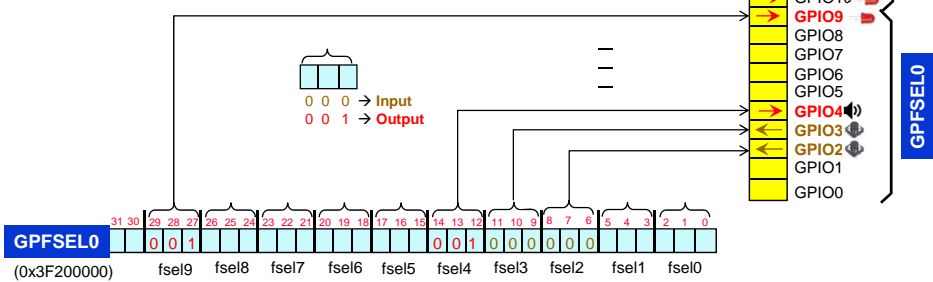
- 000: Input pin
- 001: Output pin
- 010-111: Other modes

## Configure GPIO pins as input/output

```
.set GPBASE, 0x3F200000
.set GPSEL0, 0x00
.set GPSEL1, 0x04
.set GPSEL2, 0x08
```

**Programming GPIO9 & 4 as output**  
**Programming GPIO2 & 3 as input**

```
ldr r0, =GPBASE
ldr r1, =0b000010...0001000000000000
str r1, [r0, #GPFSSEL0]
```

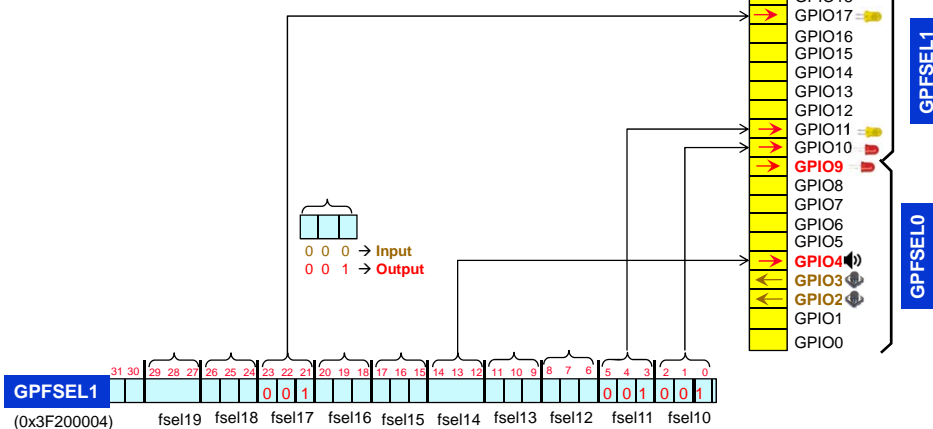


## Configure GPIO pins as input/output

**Programming GPIO10, 11 & 17 as output**

```
ldr r0, =GPBASE
ldr r1, =0x09
str r1, [r0, #GPFSSEL1]
```

```
.set GPBASE, 0x3F200000
.set GPSEL0, 0x00
.set GPSEL1, 0x04
.set GPSEL2, 0x08
```



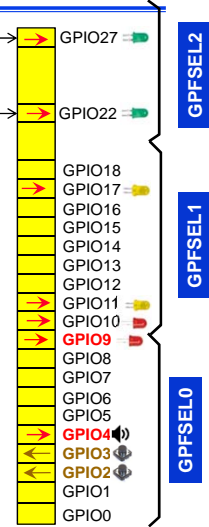
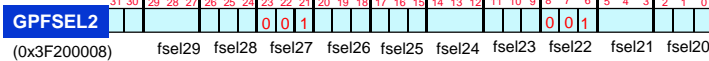
## Configure GPIO pins as input/output

### Programming GPIO22 & 27 as output

```
ldr r0, =GPBASE
ldr r1, =0x09
str r1, [r0, #GPFSEL2]
```

```
.set GPBASE, 0x3F200000
.set GPSEL0, 0x00
.set GPSEL1, 0x04
.set GPSEL2, 0x08
```

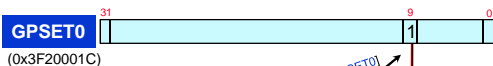
0 0 0 → Input  
0 0 1 → Output



Chapter 4.24

Dept. of Comp. Arch., UMA, 2018

## Turning ON the red LED on GPIO9



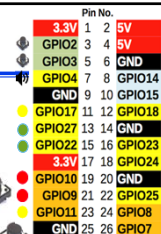
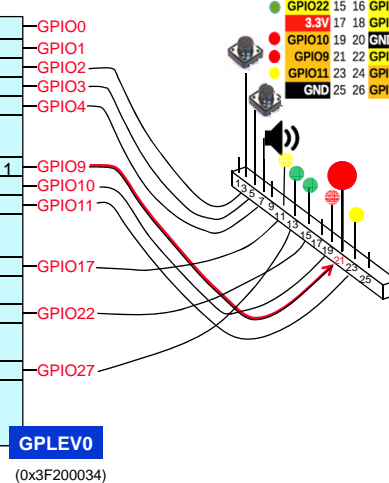
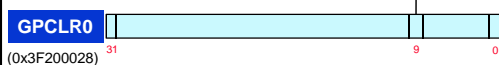
```
ldr r0, =GPBASE
ldr r1, =0b0...01000000000
str r1, [r0, #GPSET0]
```

### Turn ON Red LED (GPIO9)

```
ldr r0, =GPBASE
ldr r1, =0b0...01000000000
str r1, [r0, #GPSET0]
```

```
.set GPBASE, 0x3F200000
.set GPSET0, 0x01C
.set GPCLR0, 0x028
.set GPLEV0, 0x034
```

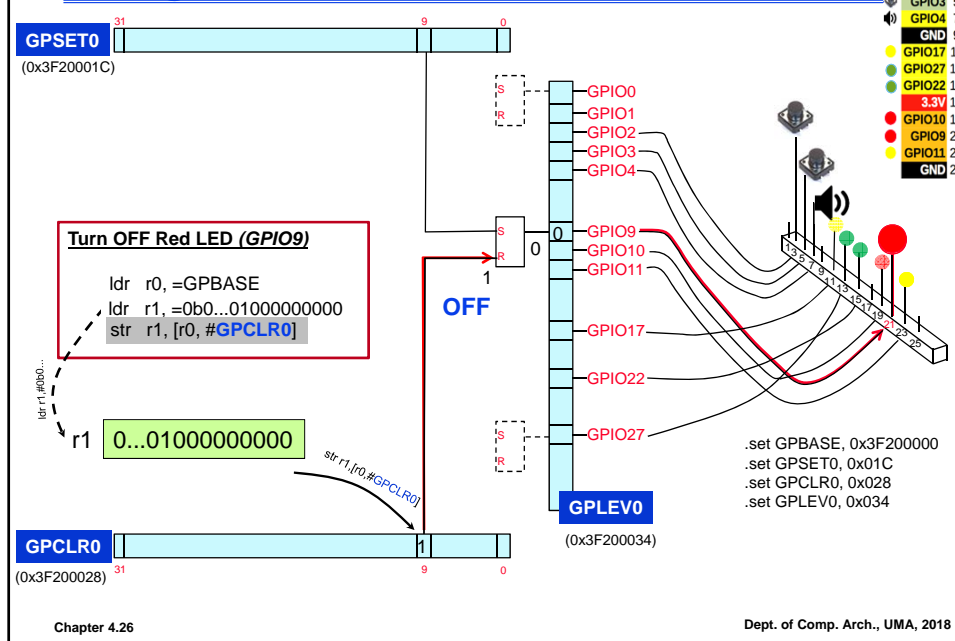
```
str r1, [r0, #GPSET0]
```



Chapter 4.25

Dept. of Comp. Arch., UMA, 2018

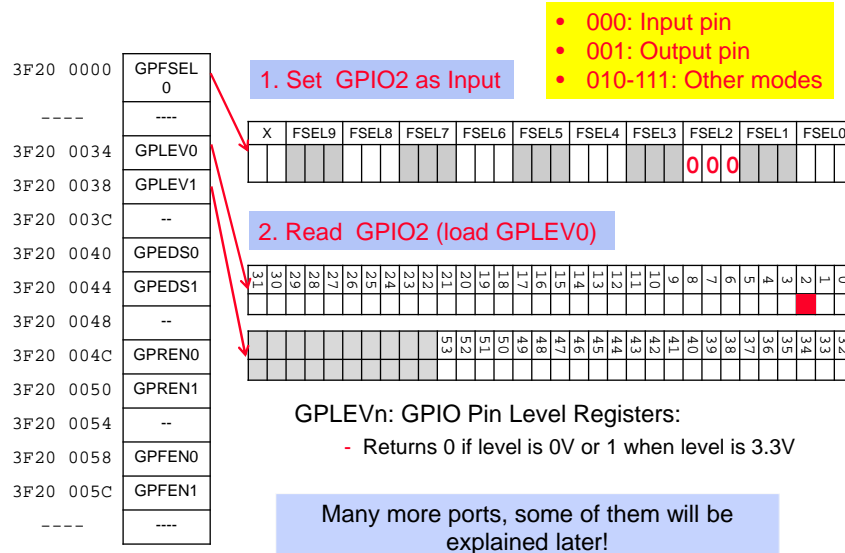
Pin	Signal
1	3.3V
3	GPI02
5	GPI03
7	GPI04
9	GND
11	GPI017
13	GPI027
15	GPI022
17	3.3V
19	GPI010
21	GPI09
23	GPI011
25	GND



❑ Code for turning on a red LED (GPIO9):

[illegible]

## GPIO memory mapping cont



## Communication of I/O Devices and Processor

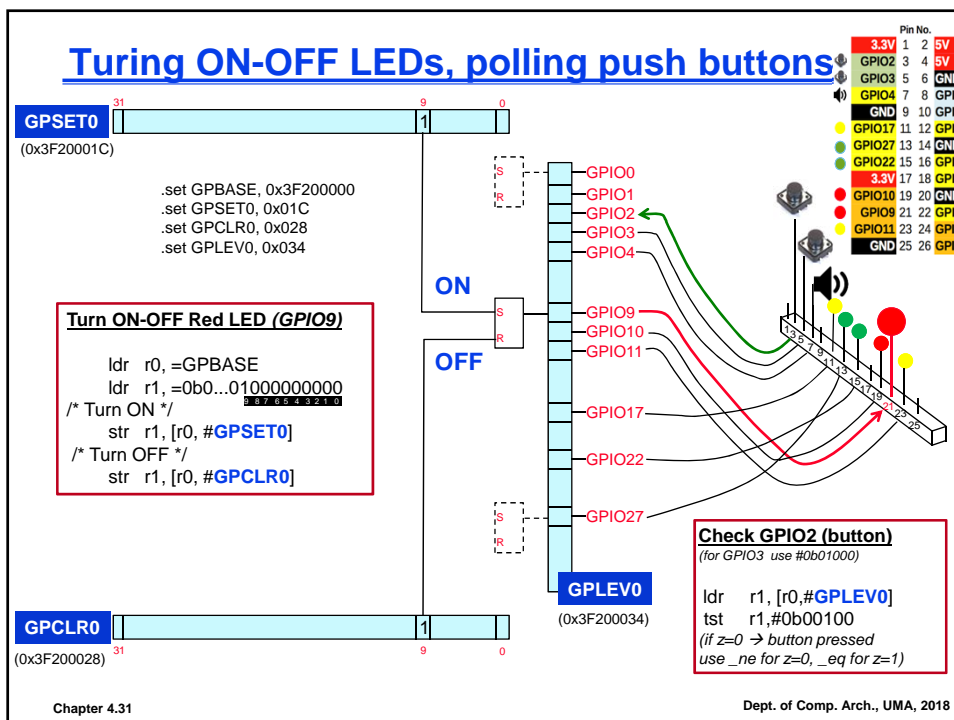
### □ How I/O devices communicate with the processor

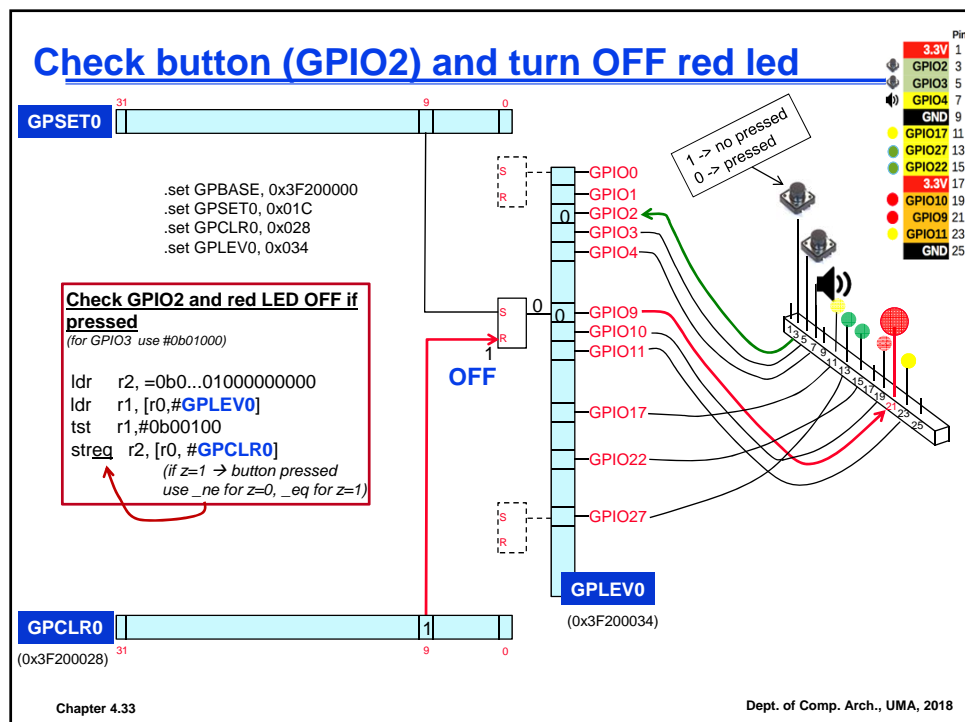
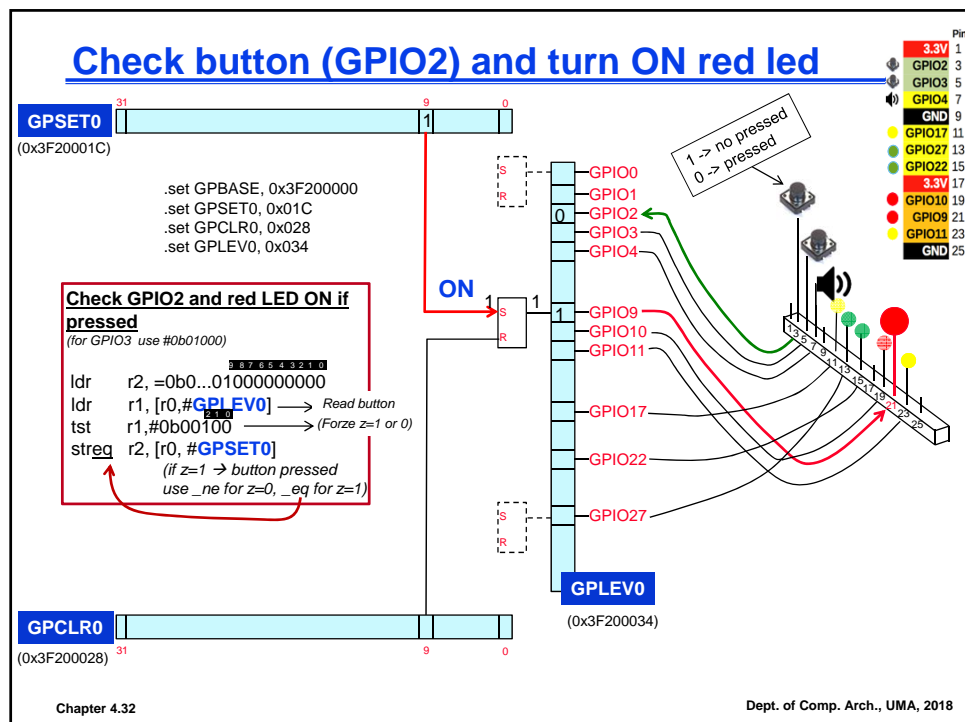
- **Polling** – the processor periodically checks the status of an I/O device (through the OS) to determine its need for service
  - Processor is totally in control – but does **all** the work
  - In real-time embedded applications:
    - I/O rates are predetermined and it makes I/O overhead predictable (helpful for real time)
  - Can waste a lot of processor time due to speed differences
- **Interrupt-driven I/O** – the I/O device issues an interrupt to indicate that it needs attention
  - Advantages of using interrupts
    - Relieves the processor from having to continuously poll for an I/O event; user program progress is only suspended during the actual transfer of I/O data to/from user memory space
  - Disadvantage – special hardware is needed to
    - Indicate the I/O device causing the interrupt and to save the necessary information prior to servicing the interrupt and to resume normal processing after servicing the interrupt

## Polling

- ❑ Periodically check I/O status register
  - If device ready, do operation
  - If error, take action
- ❑ Common in small or low-performance real-time embedded systems
  - Predictable timing
  - Low hardware cost
- ❑ In other systems, wastes CPU time

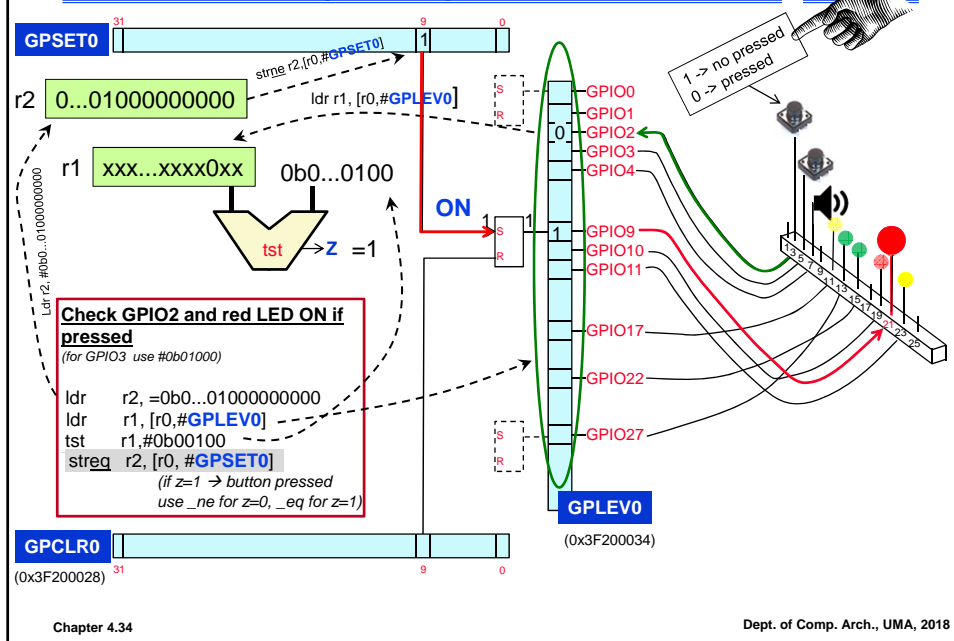
## Turing ON-OFF LEDs, polling push buttons







## Check button (GPIO2) and turn on red led



## Example 2

- Code for checking push button (GPIO2) and turning led on:

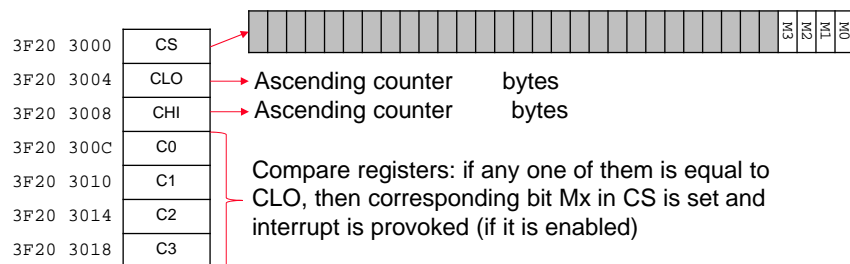
```

.set GPBASE, 0x3F200000
.set GPFSEL0, 0x00
.set GPSET0, 0x1c
.set GPLEV0, 0x34

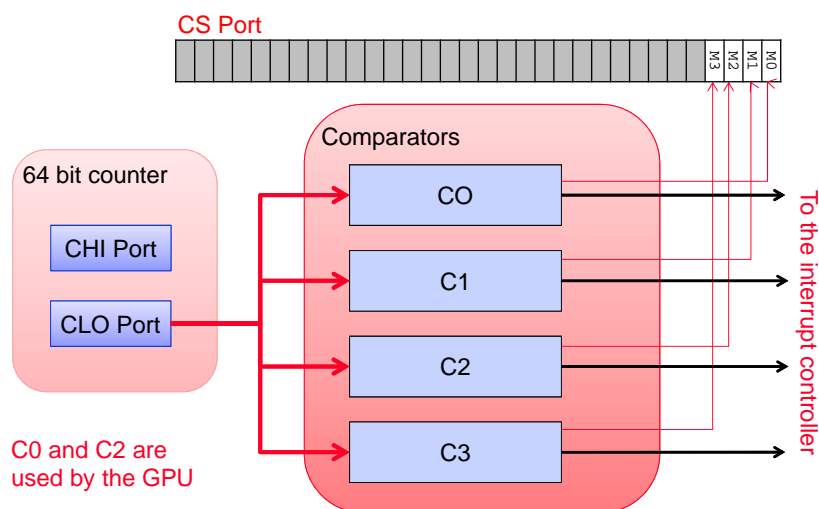
.text
ldr r0, =GPBASE
/* guia bits 10987654321098765432109876543210 */
mov r1, #0b00000000000000000000000000000000
str r1, [r0, #GPFSEL0]
/* mask for testing GPIO2 */
mov r2, #0b00000000000000000000000000000000
bucle:
ldr r3, [r0, #GPLEV0]
tst r3, r2
bne bucle
/* guia bits 10987654321098765432109876543210 */
mov r1, #0b00000000000000000000000000000000
str r1, [r0, #GPSET0]
infi: b infi
    
```

## Raspberry Pi: System Timer

- ❑ 64 bits counter (CHI:CLO → high y low)
- ❑ C0 to C3: Compare registers (4 time channels)
- ❑ CS: Control/status register
  - M0 - M3 fields are set to 1 if CLO == C0 : C3
- ❑ CK frequency: 1MHz (each increment 1 microsecond)



## System Timer



### Example 3: red LED blinking

❑ We must:

1. Configure GPIO9
2. Turn the led on
3. Wait some time
4. Turn the led off
5. Wait some time
6. Repeat steps 2-5 forever

❑ We need:

- A routine that “waits”

### Example 3: red LED blinking

```
.set GPBASE, 0x3F200000
.set GPFSEL0, 0x00
.set GPSET0, 0x1c
.set GPCLR0, 0x28
.set STBASE, 0x3F003000
.set STCLO, 0x04
```

GPIO9 configuration,  
and turning LED on and  
off

Timer read

Our “waiting” routine will:

1. Read waiting time (input parameter)
2. Repeat  
    Read current timer value  
    while it is lower than waiting time

## Example 3: red LED blinking

```
.set GPBASE, 0x3F200000
.set GPFSEL0, 0x00
.set GPSET0, 0x1c
.set GPCLR0, 0x28
.set STBASE, 0x3F003000
.set STCLO, 0x04
```

### Routine implementation:

- We can use registers r0 and r1 as input parameters
  - r0 contains the timer port address
  - r1 contains the waiting time
- We must preserve registers modified inside our routine
  - r4 contains the ending time
  - r5 loads current timer value

```
espera: push {r4, r5}           @ Save r4 and r5 in the stack
        ldr r4, [r0, #STCLO]    @ Load CLO timer
        add r4, r1              @ Add waiting time -> this is our ending
time
ret1:   ldr r5, [r0, #STCLO]    @ Enter waiting loop: load current CLO
timer   cmp r5, r4              @ Compare current time with ending time
        blo ret1               @ If lower, go back to read timer again
        pop {r4, r5}           @ Restore r4 and r5
        bx lr                  @ Return from routine
```

## Example 3: red LED blinking

```
.set GPBASE, 0x3F200000
.set GPFSEL0, 0x00
.set GPSET0, 0x1c
.set GPCLR0, 0x28
.set STBASE, 0x3F003000
.set STCLO, 0x04
```

### Our main program must:

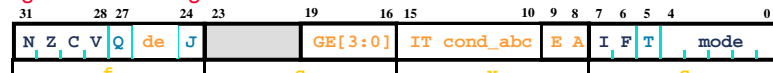
- Init the stack
- Configure GPIO9
- Init timer access (r0) and waiting time (r1) parameters
- Turn the led on and off, and call "waiting" routine between them

```
espera: push {r4, r5}           @ Save r4 and r5 in the stack
        ldr r4, [r0, #STCLO]    @ Load CLO timer
        add r4, r1              @ Add waiting time -> this is our ending
time
ret1:   ldr r5, [r0, #STCLO]    @ Enter waiting loop: load current CLO
timer   cmp r5, r4              @ Compare current time with ending time
        blo ret1               @ If lower, go back to read timer again
        pop {r4, r5}           @ Restore r4 and r5
        bx lr                  @ Return from routine
```

## Example 3: red LED blinking

```
.set GPBASE, 0x3F200000
.set GPFSEL0, 0x00
.set GPSET0, 0x1c
.set GPCLR0, 0x28
.set STBASE, 0x3F003000
.set STCLO, 0x04
```

### Program Status Register



```
espera: push {r4, r5}          @ Save r4 and r5 in the stack
        ldr r4, [r0, #STCLO]    @ Load CLO timer
        add r4, r1              @ Add waiting time -> this is our ending
time
ret1:   ldr r5, [r0, #STCLO]    @ Enter waiting loop: load current CLO
timer
        cmp r5, r4              @ Compare current time with ending time
        blo ret1                @ If lower, go back to read timer again
        pop {r4, r5}            @ Restore r4 and r5
        bx lr                    @ Return from routine
```

## Example 3: red LED blinking

```
.set GPBASE, 0x3F200000
.set GPFSEL0, 0x00
.set GPSET0, 0x1c
.set GPCLR0, 0x28
.set STBASE, 0x3F003000
.set STCLO, 0x04

.text
mov r0, #0b11010011
msr cpsr_c, r0          @ SVC mode enabled
mov sp, #0x08000000     @ Init stack in SVC mode
ldr r4, =GPBASE
mov r5, #0b00001000000000000000000000000000
str r5, [r4, #GPFSEL0]  @ Configure GPIO9
mov r5, #0b00000000000000000000000000000000
ldr r0, =STBASE         @ r0 is an input parameter (ST base address)
ldr r1, =500000         @ r1 is an input parameter (waiting time in microseconds)
bucle: bl espera        @ Call waiting routine
        str r5, [r4, #GPSET0] @ Turn LED on
        bl espera        @ Call waiting routine
        str r5, [r4, #GPCLR0] @ Turn LED off
        b bucle
espera: push {r4, r5}     @ Save r4 and r5 in the stack
        ldr r4, [r0, #STCLO] @ Load CLO timer
        add r4, r1        @ Add waiting time -> this is our ending time
ret1:   ldr r5, [r0, #STCLO] @ Enter waiting loop: load current CLO timer
        cmp r5, r4        @ Compare current time with ending time
        blo ret1          @ If lower, go back to read timer again
        pop {r4, r5}      @ Restore r4 and r5
        bx lr              @ Return from routine
```



---

# Exceptions

## Exceptions

---

- ❑ “Unexpected” events requiring change in flow of instructions execution
  - Branch and Jumps are excluded (they are “expected changes”)
- ❑ Two possible sources of exceptions
  - Internal exceptions
    - e.g., undefined opcode, overflow, syscall, ...
  - External exceptions → INTERRUPTS
    - From an external device (no memory)
- ❑ Dealing with them without sacrificing performance is hard

## Dealing with Exceptions

- ❑ Different ISAs use the terms differently
  - Traps, exceptions, interrupts ...
    - i.e.: intel x86: exceptions and interrupt
- ❑ Convention:
  - Exception: any event (other than branches and jumps) that *changes the normal flow* of instructions
    - If it is an external event the exception is called **Interrupt**

## Dealing with Exceptions

- ❑ Exceptions are just another form of control hazard. Exceptions (Interrupts) arise from
  - Arithmetic overflow (internal, exc.)
  - Trying to execute an undefined instruction (internal, exc.)
  - An OS service request (e.g., a page fault) (internal, exc.)
  - A hardware malfunction (internal or external)
  - An I/O device request (external, int)
- ❑ Invoke the OS from the user program (internal, software int. or system call)
- ❑ The software (OS) /HW looks at the cause of the exception and “deals” with it

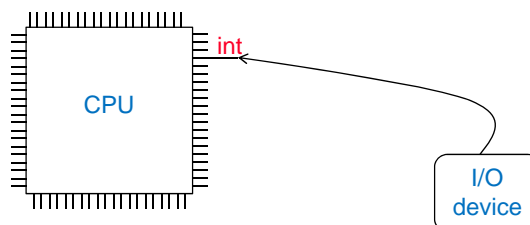


## Two Types of Exceptions

- ❑ Internal exception – synchronous to program execution
  - caused by **internal events**
  - condition must be remedied by the trap handler:
    - stop the offending instruction *midstream* in the pipeline
    - pass control to the OS trap handler
  - the offending instruction may be retried (or simulated by the OS) and the program may continue or it may be aborted

## Two Types of Exceptions

- ❑ External exceptions -> **Interrupts** – asynchronous to program execution
  - caused by **external events**
  - may be handled **between** instructions:
    - let the prior instructions currently active in the pipeline *complete*
    - pass control to the OS interrupt handler
  - simply suspend and resume user program



## Interrupt Driven I/O

- ❑ An I/O interrupt is **asynchronous** wrt instruction execution
  - Is not associated with any instruction so doesn't prevent any instruction from completing
    - You can pick your own convenient point to handle the interrupt
    - Control unit needs only check for a pending I/O interrupt at the time it starts a new instruction
- ❑ With I/O interrupts
  - Need a way to identify the device generating the interrupt
    - **Vectored interrupts**: the device can send a vector (id.) to the processor, which uses it to address the table of the interrupt vectors, from where it gets the address of the handle.
    - **Non vectored interrupts**: the device places a status field in the Cause register, jumps to a handler at a fixed direction.
    - **Auto-vectored interrupts**: each exception has vector associated to it.
    - When the handle gets control, it knows the identity of the device and can immediately start the I/O operation
  - Can have different urgencies (so need a way to **prioritize** them)
    - I/O interrupts have lower priority than internal exceptions
    - UNIX OS uses four to six levels
  - ❑ Interrupt priority levels (IPLs) assigned by the OS to each process can be raised and lowered via changes to the Status's Interrupt mask field
    - Lowest IPL: all interrupts are permitted
    - Highest IPL: all interrupts are blocked

## Exceptions in ARM

- ❑ ARM's exception system is auto-vectorized
  - There are 8 exception types, NI=0:7
  - Each NI has an exception vector associated to it
    - The exception vector is a jump to a handler
    - $NI \cdot 4$  is the offset to the exception vectors table

Exception	Type	Offset	Mode
Reset	Interrupt	0x00	SVC
Undefined Instruct.	Exception	0x04	Undefined
SW interrupt	SW Interrup.	0x08	SVC
Prefetch abort	Exception	0x0C	Abort
Data abort	Exception	0x10	Abort
Reserved	-	0x14	-
IRQ	Interrupt	0x18	IRQ
FIQ	Interrupt	0x1C	FIQ

## Exceptions in ARM

### □ Type of exceptions:

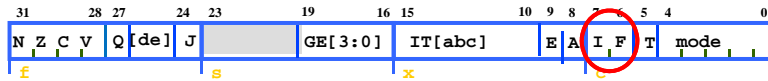
- Reset: pins in P6 fire a bootloader
- Undefined instruction: op. code not valid
- Software interruptions: system calls
- Prefetch abort /data abort: memory misalignment, access privilege errors
- IRQ: interruptions due to external devices
- FIQ: fast interruptions

## Exception priorities

- When multiple exceptions arise at the same time, a fixed priority system determines the order that they are handled:

Priority		Exception
Higuest	1	Reset
	2	Precise Data Abort
	3	FIQ
	4	IRQ
	5	Prefetch Abort
Lowest	6	Imprecise Data Abort
	7	BKPT
		Undefined Instruction
		SVC SMC

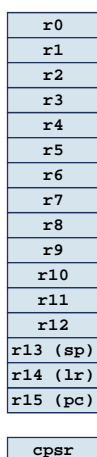
## Status register again: cpsr fsxc



- **Condition code flags**
  - N = **N**egative result from ALU
  - Z = **Z**ero result from ALU
  - C = ALU operation **C**arried out
  - V = ALU operation **o**verflowed
- **Sticky Overflow flag - Q flag**
  - Indicates if saturation has occurred
- **SIMD Condition code bits – GE[3:0]**
  - Used by some SIMD instructions
- **IF THEN status bits – IT[abcde]**
  - Controls conditional execution of Thumb instructions
- **T bit**
  - T = 0: Processor in ARM state
  - T = 1: Processor in Thumb state
- **J bit**
  - J = 1: Processor in Jazelle state
- **Mode bits**
  - Specify the processor mode
- **Interrupt Disable bits**
  - I = 1: Disables IRQ
  - F = 1: Disables FIQ
- **E bit**
  - E = 0: Data load/store is little endian
  - E = 1: Data load/store is bigendian
- **A bit**
  - A = 1: Disable imprecise data aborts

## Banking of registers

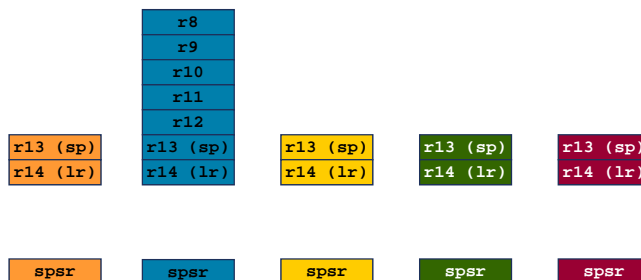
10000      10010    10001    11011    10111    10011  
**User mode**      **IRQ**      **FIQ**      **Undef**      **Abort**      **SVC**



**Current mode**

- **ARM has 37 registers, all 32-bits long**

- A subset of these registers is accessible in each mode and does not have to be preserved  
 - Note: System mode uses the User mode register set.



**Banked out registers**

## Handling exceptions

- **When an exception occurs, the core...**
  - Copies CPSR into SPSR\_<mode>
  - Sets appropriate CPSR bits
    - Change to ARM state
    - Change to exception mode
    - Disable interrupts (if appropriate)
  - Stores the return address in LR\_<mode>
  - Sets PC to vector address
- **To return, exception handler needs to...**
  - Restore CPSR from SPSR\_<mode>
  - Restore PC from LR\_<mode>

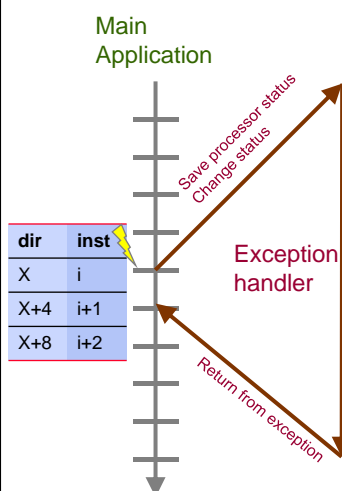
	...
0x1C	FIQ
0x18	IRQ B irq_handler
0x14	(Reserved)
0x10	Data Abort
0x0C	Prefetch Abort
0x08	Supervisor Call
0x04	Undefined Instruction
0x00	Reset

**Vector Table**

Vector table can also be at  
0xFFFF0000 on most cores

This can only be done in ARM state.

## Handling exceptions



### 1. Save processor status

- Stores PC in LR\_<mode>
  - Adjusts LR based on exception type
  - Stores X+8 → LR\_<mode>
- Copies CPSR into SPSR\_<mode>

### 2. Change processor status for exception

- Forces the CPSR mode bits to a value (depends on the exception)
- Sets PC to vector address

### 3. Execute exception handler

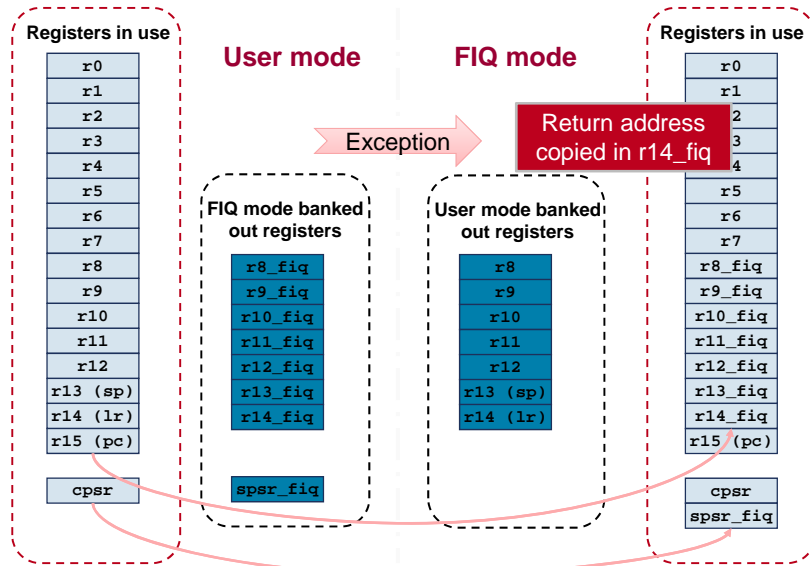
- <user code>

### 4. Return to main application

- Restore CPSR from SPSR\_<mode>
- Restore PC: PC ← LR\_<mode> - 4

- 1 and 2 performed automatically by the core
- 3 and 4 responsibility of software

## User mode → FIQ mode



Chapter 4.61

Dept. of Comp. Arch., UMA, 2018

## Exception handler

### □ Basic structure of a exception handler

- Interruption: the return is done by lr-4
- Internal exception (as data abort): the return is done by lr-8
- User must manage A, I and F flags to disable/enable nesting of new exceptions and interruptions.
  - Initially the interruptions are disabled (I=F=1).

**irq\_handler:**

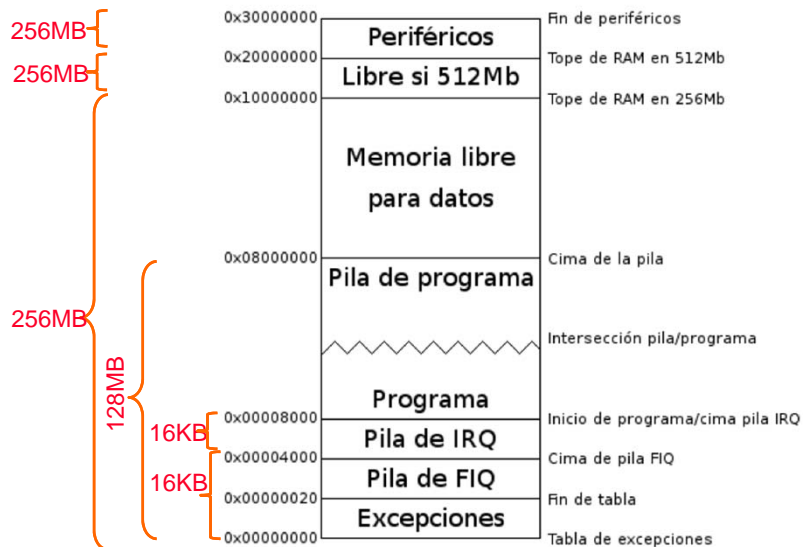
- ① Push registers to be used
- ② Source of interruption?
- ③ Perform handler work depending on ②
- ④ Clear event (notify to device IRQ/FIQ has been served)
- ⑤ Pop registers
- ⑥ Return from handler: `subs pc, lr, #4`

Chapter 4.62

Dept. of Comp. Arch., UMA, 2018

## Memory map

- Exception Vector Table starts at 0x00000000 (or 0xFFFF0000)



## Main program: steps to set up the Interrupts

- 1 Initialize Vector Table (IRQ/FIQ) in the Vector Table
- 2 Init the stack/s for FIQ/IRQ modes
 

```
sp_ifq <- 0x00004000
sp_irq <- 0x00008000
```
- 3 Init the stack for SVC mode (SVC mode selected)
 

```
sp_svc <- 0x08000000
```
- 4 Configure GPIOs (I&O)
- 5 Configure peripheral interruption: timer/push-buttons
- 6 Local enabling of configured interrupts
- 7 Global enabling of interrupts (SVC mode)
- 8 ... (main program tasks)

## ① Initialize Vector Table

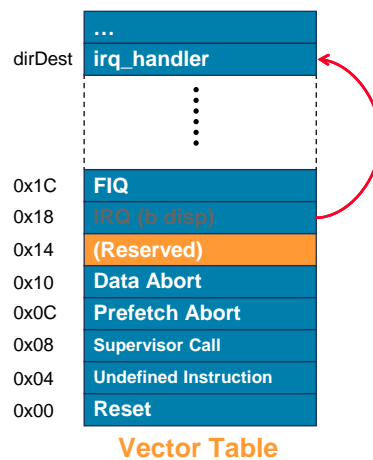
- ❑ To write in the Vector Table we can use a macro, ADDEXC, that computes the offset of the exception handler and writes the Vector in the Vector Table.

```
mov r0, #0 @Vector table base = 0
ADDEXC 0x18, irq_handler
```

## macro ADDEXC offset, dirDest

- ❑ The IRQ handler is located at `dirDest`
- ❑ Vector table stores a branch instruction to the IRQ handler, `b disp`, located at offset (0x18)
- ❑ `disp` is the number of bytes between `dirDest` and offset, divided by 4
- ❑ While executing `b disp`, `pc` is incremented twice (`pc = 0x18+8`)
- ❑ Thus, `disp` must store

$$disp = \frac{(dirDest - (offset + 8))}{4}$$





## macro ADDEXC offset, dirDest

ASM	b	disp
Binary	11101010	
Hex	E A	

$$disp = \frac{(dirDest - (offset + 8))}{4}$$

$$0xEA000000 + \frac{(dirDest - (offset + 8))}{4} =$$

$$0xE9FFFFFFE + \frac{(dirDest - offset)}{4} =$$

8 digits!

$$0x3A7FFFFFF8 + \frac{(dirDest - offset)}{4}$$

We can divide by 4 by rotating right the numerator (after adding 3 to save the two most significant bits)

dirDest	...
0x1C	irq_handler
0x18	...
0x14	FIQ
0x10	IRQ (b-disp)
0x0C	(Reserved)
0x08	Data Abort
0x04	Prefetch Abort
0x00	Supervisor Call
	Undefined Instruction
	Reset

**Vector Table**

## macro ADDEXC offset, dirDest

ASM	b	disp
Binary	11101010	
Hex	E A	

$$disp = \frac{(dirDest - (offset + 8))}{4}$$

```

macro ADDEXC offset, dirDest
    ldr r1, =(\dirDest-\offset+0xA7FFFFFFB)
    ror r1, #2
    str r1, [r0, #\offset]
.endm
    
```

$$0xEA000000 + \frac{(dirDest - (offset + 8))}{4} =$$

$$0xE9FFFFFFE + \frac{(dirDest - offset)}{4} =$$

$$0x3A7FFFFFF8 + \frac{(dirDest - offset)}{4}$$

We can divide by 4 by rotating right the numerator (after adding 3 to save the two most significant bits)

dirDest	...
0x1C	irq_handler
0x18	...
0x14	FIQ
0x10	IRQ (b-disp)
0x0C	(Reserved)
0x08	Data Abort
0x04	Prefetch Abort
0x00	Supervisor Call
	Undefined Instruction
	Reset

**Vector Table**

## ②③ Initialize the stack

- ❑ Each mode has its stack pointer (sp)
  - Change the mode (via `cpsr_c`)
  - Instructions `msr (sr <- reg)` and `mrs (reg <- sr)`.
  - Initialize the corresponding sp register
- ❑ Initial state in BareMetal is SVC
  - `sp_fiq=0x4000, sp_irq=0x8000, sp_svc=0x08000000:`

```
mov     r0, #0 @ Pointer to vector table
ADDEXC 0x18, irq_handler
ADDEXC 0x1c, fiq_handler
```

```
2  msr    cpsr_c, r0
    mov    sp, #0x4000
```

```

mov    r0, #0b11010010 @ IRQ mode, FIQ and IRQ disabled
msr    cpsr_c, r0
mov    sp, #0x8000

```

```

mov    r0, #0b11010011 @ SVC mode, FIQ and IRQ disabled
3 msr    cpsr_c, r0
mov    sp, #0x08000000

```

## 5 Configure peripheral interruption

- ❑ GPIO interruption (push-buttons): use GPRENn, GPFENn, GPHENn, GPLENn, GPARENn y GPAFENn

### 5 Configure push-button interruption (GPIO 2)

[illegible]

- ❑ **System Timer:** write the final count (microseconds) in STC1/STC3

### 5 Configure timer IRQ

```
ldr r0, =STBASE
ldr r1, [r0, #STCLO]
add r1, #y @y microseconds
str r1, [r0, #STC1]
```

## GPIO memory mapping

## Chapter 4.71

High ENable → Enable interrupt request when pin has 1

Low ENable → Enable interrupt request when pin has 0

Async. Rising edge ENable → Enable  
interrupt request with a rising edge (async.,  
glitch possible)

Async. Falling edge ENable → Enable  
interrupt request with a falling edge (async.,  
glitch possible)

### Pull Up Down control

### Pull Up Down clock

## 2 Source of interruption?

- ❑ In case of interruption, the handler must identify the source reading the IRQ pending ports

- GPIO interruption detection: use GPEDSn.

2 Source of interruption?. Check if push-button (2) was pressed

[illegible]

- System Timer interrupt detection: STCS notifies interruption due to C0 : C3 counters

2 Source of timer interruption?:

```
ldr r0, #STBASE
ldr r2, [r0, #STCS]
ands r2, #0b0010 @C1?
...
ldr r2, [r0, #STCS]
ands r2, #0b1000 @C3?
```

## GPIO memory mapping

3F20 0000	GPFSSEL0
----	----
3F20 0034	GPLEV0
3F20 0038	GPLEV1
3F20 003C	--
3F20 0040	GPEDS0
3F20 0044	GPEDS1
3F20 0048	--
3F20 004C	GPREN0
3F20 0050	GPREN1
3F20 0054	--
3F20 0058	GPFEN0
3F20 005C	GPFEN1
----	----

Event Detect Status→ Set 1 when interrupt is requested. Must be cleared writing 1 when the interrupt has been serviced

Rising edge ENable→ Enable interrupt request with a rising edge (sync., avoid glitch)

Falling edge ENable→ Enable interrupt request with a falling edge (sync., avoid glitch)

The suppression of glitches is done by sampling the pin using the system clock and then looking for a "011" (rising) or "011" (falling edge) pattern on the sampled signal.

## ⑥ Local enabling of sources of interruption

3F00 B200	IRQ basic pending	ARM client	System_G0	0	0	ARM mailbox	ARM mailbox
3F00 B204	IRQ pending 1	ARM mailbox	System_G1	1	2	ARM mailbox	ARM mailbox
3F00 B208	IRQ pending 2	ARM mailbox	System_G2	2	3	ARM mailbox	ARM mailbox
3F00 B20C	FIQ control	ARM mailbox	System_G3	3	4	ARM mailbox	ARM mailbox
3F00 B210	Enable IRQs 1	ARM mailbox	System_G4	4	5	ARM mailbox	ARM mailbox
3F00 B214	Enable IRQs 2	ARM mailbox	System_G5	5	6	ARM mailbox	ARM mailbox
3F00 B218	Enable Basic IRQs	ARM mailbox	System_G6	6	7	ARM mailbox	ARM mailbox
3F00 B21C	Disable IRQs 1	ARM mailbox	System_G7	7	8	ARM mailbox	ARM mailbox
3F00 B220	Disable IRQs 2	ARM mailbox	System_G8	8	9	ARM mailbox	ARM mailbox
3F00 B224	Disable Basic IRQs	ARM mailbox	System_G9	9	10	ARM mailbox	ARM mailbox

## ⑥ Local enabling of sources of interruption

- Three groups: pending, enable and disable

3F00 B200	IRQ basic pending	Pending
3F00 B204	IRQ pending 1	
3F00 B208	IRQ pending 2	
3F00 B20C	FIQ control	Enable
3F00 B210	Enable IRQs 1	
3F00 B214	Enable IRQs 2	
3F00 B218	Enable Basic IRQs	Disable
3F00 B21C	Disable IRQs 1	
3F00 B220	Disable IRQs 2	
3F00 B224	Disable Basic IRQs	

In each group:

- IRQ basic: summary
- IRQs 1 and 2: in detail

There is also one port for FIQ control

## ⑥ Local enabling of sources of interruption

FIQ Control		ARM timer	31
		ARM MailBox	30
		ARM MailBox 0	29
		ARM MailBox 1	28
		cpu0 halted	27
		cpu1 halted	26
		illegal type 0	25
		illegal type 1	24
		illegal type 2	23
		illegal type 3	22
		illegal type 4	21
		illegal type 5	20
		illegal type 6	19
		illegal type 7	18
		illegal type 8	17
		illegal type 9	16
		illegal type 10	15
		illegal type 11	14
		illegal type 12	13
		illegal type 13	12
		illegal type 14	11
		illegal type 15	10
		illegal type 16	9
		illegal type 17	8
		illegal type 18	7
		illegal type 19	6
		illegal type 20	5
		illegal type 21	4
		illegal type 22	3
		illegal type 23	2
		illegal type 24	1
		illegal type 25	0
		system_c1	31
		system_c2	30
		system_c3	29
		system_c4	28
		system_c5	27
		system_c6	26
		system_c7	25
		system_c8	24
		system_c9	23
		system_c10	22
		system_c11	21
		system_c12	20
		system_c13	19
		system_c14	18
		system_c15	17
		system_c16	16
		system_c17	15
		system_c18	14
		system_c19	13
		system_c20	12
		system_c21	11
		system_c22	10
		system_c23	9
		system_c24	8
		system_c25	7
		system_c26	6
		system_c27	5
		system_c28	4
		system_c29	3
		system_c30	2
		system_c31	1
		system_c32	0
		system_c33	31
		system_c34	30
		system_c35	29
		system_c36	28
		system_c37	27
		system_c38	26
		system_c39	25
		system_c40	24
		system_c41	23
		system_c42	22
		system_c43	21
		system_c44	20
		system_c45	19
		system_c46	18
		system_c47	17
		system_c48	16
		system_c49	15
		system_c50	14
		system_c51	13
		system_c52	12
		system_c53	11
		system_c54	10
		system_c55	9
		system_c56	8
		system_c57	7
		system_c58	6
		system_c59	5
		system_c60	4
		system_c61	3
		system_c62	2
		system_c63	1
		system_c64	0
		system_c65	31
		system_c66	30
		system_c67	29
		system_c68	28
		system_c69	27
		system_c70	26
		system_c71	25
		system_c72	24
		system_c73	23
		system_c74	22
		system_c75	21
		system_c76	20
		system_c77	19
		system_c78	18
		system_c79	17
		system_c80	16
		system_c81	15
		system_c82	14
		system_c83	13
		system_c84	12
		system_c85	11
		system_c86	10
		system_c87	9
		system_c88	8
		system_c89	7
		system_c90	6
		system_c91	5
		system_c92	4
		system_c93	3
		system_c94	2
		system_c95	1
		system_c96	0
		system_c97	31
		system_c98	30
		system_c99	29
		system_c100	28
		system_c101	27
		system_c102	26
		system_c103	25
		system_c104	24
		system_c105	23
		system_c106	22
		system_c107	21
		system_c108	20
		system_c109	19
		system_c110	18
		system_c111	17
		system_c112	16
		system_c113	15
		system_c114	14
		system_c115	13
		system_c116	12
		system_c117	11
		system_c118	10
		system_c119	9
		system_c120	8
		system_c121	7
		system_c122	6
		system_c123	5
		system_c124	4
		system_c125	3
		system_c126	2
		system_c127	1
		system_c128	0
		system_c129	31
		system_c130	30
		system_c131	29
		system_c132	28
		system_c133	27
		system_c134	26
		system_c135	25
		system_c136	24
		system_c137	23
		system_c138	22
		system_c139	21
		system_c140	20
		system_c141	19
		system_c142	18
		system_c143	17
		system_c144	16
		system_c145	15
		system_c146	14
		system_c147	13
		system_c148	12
		system_c149	11
		system_c150	10
		system_c151	9
		system_c152	8
		system_c153	7
		system_c154	6
		system_c155	5
		system_c156	4
		system_c157	3
		system_c158	2
		system_c159	1
		system_c160	0
		system_c161	31
		system_c162	30
		system_c163	29
		system_c164	28
		system_c165	27
		system_c166	26
		system_c167	25
		system_c168	24
		system_c169	23
		system_c170	22
		system_c171	21
		system_c172	20
		system_c173	19
		system_c174	18
		system_c175	17
		system_c176	16
		system_c177	15
		system_c178	14
		system_c179	13
		system_c180	12
		system_c181	11
		system_c182	10
		system_c183	9
		system_c184	8
		system_c185	7
		system_c186	6
		system_c187	5
		system_c188	4
		system_c189	3
		system_c190	2
		system_c191	1
		system_c192	0
		system_c193	31
		system_c194	30
		system_c195	29
		system_c196	28
		system_c197	27
		system_c198	26
		system_c199	25
		system_c200	24
		system_c201	23
		system_c202	22
		system_c203	21
		system_c204	20
		system_c205	19
		system_c206	18
		system_c207	17
		system_c208	16
		system_c209	15
		system_c210	14
		system_c211	13
		system_c212	12
		system_c213	11
		system_c214	10
		system_c215	9
		system_c216	8
		system_c217	7
		system_c218	6
		system_c219	5
		system_c220	4
		system_c221	3
		system_c222	2
		system_c223	1
		system_c224	0
		system_c225	31
		system_c226	30
		system_c227	29
		system_c228	28
		system_c229	27
		system_c230	26
		system_c231	25
		system_c232	24
		system_c233	23
		system_c234	22
		system_c235	21
		system_c236	20
		system_c237	19
		system_c238	18
		system_c239	17
		system_c240	16
		system_c241	15
		system_c242	14
		system_c243	13
		system_c244	12
		system_c245	11
		system_c246	10
		system_c247	9
		system_c248	8
		system_c249	7
		system_c250	6
		system_c251	5
		system_c252	4
		system_c253	3
		system_c254	2
		system_c255	1
		system_c256	0
		system_c257	31
		system_c258	30
		system_c259	29
		system_c260	28
		system_c261	27
		system_c262	26
		system_c263	25
		system_c264	24
		system_c265	23
		system_c266	22
		system_c267	21
		system_c268	20
		system_c269	19
		system_c270	18
		system_c271	17
		system_c272	16
		system_c273	15
		system_c274	14
		system_c275	13
		system_c276	12
		system_c277	11
		system_c278	10
		system_c279	9
		system_c280	8
		system_c281	7
		system_c282	6
		system_c283	5
		system_c284	4
		system_c285	3
		system_c286	2
		system_c287	1
		system_c288	0
		system_c289	31
		system_c290	30
		system_c291	29
		system_c292	28
		system_c293	27
		system_c294	26
		system_c295	25
		system_c296	24
		system_c297	23
		system_c298	22
		system_c299	21
		system_c300	20
		system_c301	19
		system_c302	18
		system_c303	17
		system_c304	16
		system_c305	15
		system_c306	14
		system_c307	13
		system_c308	12
		system_c309	11
		system_c310	10
		system_c311	9
		system_c312	8
		system_c313	7
		system_c314	6
		system_c315	5
		system_c316	4
		system_c317	3
		system_c318	2
		system_c319	1
		system_c320	0
		system_c321	31
		system_c322	30
		system_c323	29
		system_c324	28
		system_c325	27
		system_c326	26
		system_c327	25
		system_c328	24
		system_c329	23
		system_c330	22
		system_c331	21
		system_c332	20
		system_c333	19
		system_c334	18
		system_c335	17
		system_c336	16
		system_c337	15
		system_c338	14
		system_c339	13
		system_c340	12
		system_c341	11
		system_c342	10
		system_c343	9
		system_c344	8
		system_c345	7
		system_c346	6
		system_c347	5
		system_c348	4
		system_c349	3
		system_c350	2
		system_c351	1
		system_c352	0
		system_c353	31
		system_c354	30
		system_c355	29
		system_c356	28
		system_c357	27
		system_c358	26
		system_c359	25
		system_c360	24
		system_c361	23
		system_c362	22
		system_c363	21
		system_c364	20
		system_c365	19
		system_c366	18
		system_c367	17
		system_c368	16
		system_c369	15
		system_c370	14
		system_c371	13
		system_c372	12
		system_c373	11
		system_c374	10
		system_c375	9
		system_c376	8
		system_c377	7
		system_c378	6
		system_c379	5
		system_c380	4
		system_c381	3
		system_c382	2
		system_c383	1
		system_c384	0
		system_c385	31
		system_c386	30
		system_c387	29
		system_c388	28
		system_c389	27
		system_c390	26
		system_c391	25
		system_c392	24
		system_c393	23
		system_c394	22
		system_c395	21
		system_c396	20
		system_c397	19
		system_c398	18
		system_c399	17
		system_c400	16
		system_c401	15
		system_c402	14
		system_c403	13
		system_c404	12
		system_c405	11
		system_c406	10
		system_c407	9
		system_c408	8
		system_c409	7
		system_c410	6
		system_c411	5
		system_c412	4
		system_c413	3
		system_c414	2
		system_c415	1
		system_c416	0
		system_c417	31
		system_c418	30
		system_c419	29
		system_c420</	

## Example 5: turn on a red led after 4 seconds

```
.include "inter.inc"

.text
mov r0, #0
ADDEXC 0x18, irq_handler

. . .

ldr r0, =GPFASE
ldr r1, =0b00001000000000000000000000000000
str r1, [r0, #GPFSEL0]

ldr r0, =STRASE
ldr r1, [r0, #STCLO]
add r1, #0x400000 @ 4.19 seconds
str r1, [r0, #STC1]

ldr r0, =INTBASE
mov r1, #0b0010
str r1, [r0, #INTENIRQ1]

mov r0, #0b01010011 @ SVC mode, IRQ enabled
msr cpsr_c, r0

buc: b buc

irq_handler:
push {r0, r1}

ldr r0, =GPFASE
mov r1, #0b0000000000000000000000001000000000
str r1, [r0, #GPSET0]

pop {r0, r1}
subs pc, lr, #4
```

Initialize vector table ①

Set GPIO9 as Output ④

Load CLO, add 4 sec and store result in C1 ⑤

Enable C1 interruption ⑥

Enable I flag ⑦

①

Turn on RED led (GPIO9) ③

PC ← LR - 4 ⑥

⑤

## Example 5: inter.inc file

```
.macro ADDEXC vector, dirRTI
ldr r1, =(\dirRTI-\vector+0xa7fffffb)
ror r1, #2
str r1, [r0, #\vector]
.endm

.set GPFASE, 0x3F200000
.set GPFSEL0, 0x00
.set GPFSEL1, 0x04
.set GPFSEL2, 0x08
.set GPFSEL3, 0x0c
.set GPFSEL4, 0x10
.set GPFSEL5, 0x14
.set GPFSEL6, 0x18
.set GPSET0, 0x1c
.set GPSET1, 0x20
.set GPCLR0, 0x28
.set GPCLR1, 0x2c
.set GPLEV0, 0x34
.set GPLEV1, 0x38
.set GPEDS0, 0x40
.set GPEDS1, 0x44
.set GPFEN0, 0x58
.set GPFEN1, 0x5c
.set GPPUD, 0x94
.set GPPUDCLK0, 0x98

.set STRASE, 0x3F003000
.set STCS, 0x00
.set STCLO, 0x04
.set STC1, 0x10
.set STC3, 0x18

.set INTBASE, 0x3F00b000
.set INTFIOCON, 0x20c
.set INTENIRQ1, 0x210
.set INTENIRQ2, 0x214
```

GPIO

Timer

Interrupt

## Example 6: turn on a red led after pushing a button

```

.include "inter.inc"
.text
mov     r0, #0
ADDEXC 0x18, irq_handler
1
mov     r0, #0b11010010
msr     cpsr_c, r0
2
mov     sp, #0x8000
Stack init for IRQ mode
3
mov     r0, #0b11010011
msr     cpsr_c, r0
Stack init for SVC mode
4
ldr     r0, =GPBASE
mov     r1, #0b00001000000000000000000000000000
Set GPIO9 as output
5
str     r1, [r0, #GPFSEL0]
mov     r1, #0b00000000000000000000000000000000
Configure falling edge
6
str     r1, [r0, #GPFEN0]
interruptions through GPIO2
7
ldr     r0, =INTBASE
mov     r1, #0b00000000000100000000000000000000
Allow interruptions
8
str     r1, [r0, #INTENIRQ2]
from any GPIO pin
mov     r0, #0b01010011
msr     cpsr_c, r0
Set SVC mode with IRQ enabled
buc:    b      buc

```

## Example 6: IRQ handler

```

irq_handler:
1
push    {r0, r1, r2}
ldr     r0, =GPBASE
ldr     r2, [r0, #GPFDS0]
2
ands    r2, #0b00000000000000000000000000000000
Check GPIO2 was pressed
3
movne   r1, #0b00000000000000000000000000000000
Turn on GPIO9 red led
4
strne   r1, [r0, #GPSET0]
Clear GPIO2 event
5
movne   r1, #0b00000000000000000000000000000000
6
strne   r1, [r0, #GPFDS0]
pop     {r0, r1, r2}
subs    pc, lr, #4

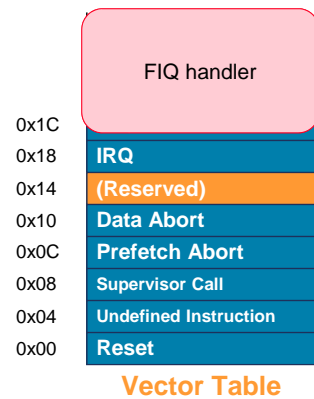
```

## Using FIQ

### Advantages:

- Registers r8 to r14 are saved
- FIQ handler can start after vector table: no need to branch

### Disadvantage: only one source of interruption can be handled



## Using FIQ

### FIQ

Select FIQ Source  
FIQ Enable

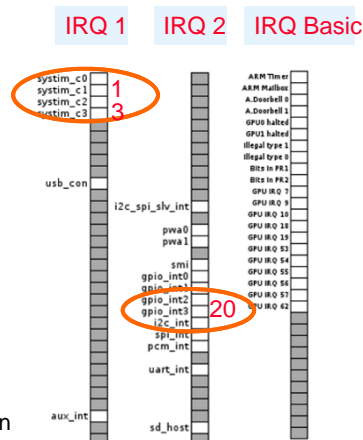
- Select FIQ Source = 7 bits → 128 sources
  - 0-31 represent 32 interruption sources of IRQ 1
  - 32-63 represent 32 interruption sources of IRQ 2
  - 64-95 represent 32 interruption sources of IRQ basic

### Examples:

- Enable FIQ for C1 of SysTimer
  - Bit 1 of IRQ1 → Code 1
  - Also 1 in FIQ Enable
  - Result: 0b10000001 → 0x81
- Enable FIQ for C3 of SysTimer
  - Bit 3 of IRQ1 → Code 3
  - Also 1 in FIQ Enable
  - Result: 0b10000011 → 0x83
- Enable FIQ for GPIO\_int3
  - Bit 20 of IRQ2 → Code 20+32
  - Also 1 in FIQ Enable
  - Result: 0b10110100 → 0xB4

### Disadvantage:

- Just one source interruption can be enabled!





## Example 7: Putting it all together

- ❑ Turn on and off the red led at GPIO9 every 4 seconds **or** when the push button at GPIO2 is pressed.
- ❑ Use IRQ to handle the timer and FIQ to handle the push button.
- ❑ Use a variable in memory to control the led state (on or off)

## Example 7: Timer in IRQ and push button in FIQ

```
.include "inter.inc"
.text
ADDEXC 0x18, irq_handler
ADDEXC 0x1c, fiq_handler

mov r0, #0b11010001
msr cpsr_c, r0
mov sp, #0x4000
Stack init for FIQ mode

mov r0, #0b11010010
msr cpsr_c, r0
mov sp, #0x8000
Stack init for IRQ mode

mov r0, #0b11010011
msr cpsr_c, r0
mov sp, #0xb000000
Stack init for SVC mode

ldr r0, =GPBASE
mov r1, #0b00001000000000000000000000000000
str r1, [r0, #GPFSEL0]
Set GPIO9 as output

mov r1, #0b00000000000000000000000000000000
str r1, [r0, #GPFEN0]
Configure FE ints through GPIO2

ldr r0, =STBASE
ldr r1, [r0, #STCLO]
add r1, #0x400000
str r1, [r0, #STC1]
Program timer to interrupt in 4 seconds

ldr r0, =INTBASE
mov r1, #0b000000010
str r1, [r0, #INTENIRQ1]
Enable C1 interruption

mov r1, #0b10110100
str r1, [r0, #INTFICON]
Enable FIQ for GPIO_int3

mov r0, #0b00010011
msr cpsr_c, r0
Set SVC mode with FIQ and IRQ enabled

bucle: b bucle
```

```

    onoff: .word 0

```

---

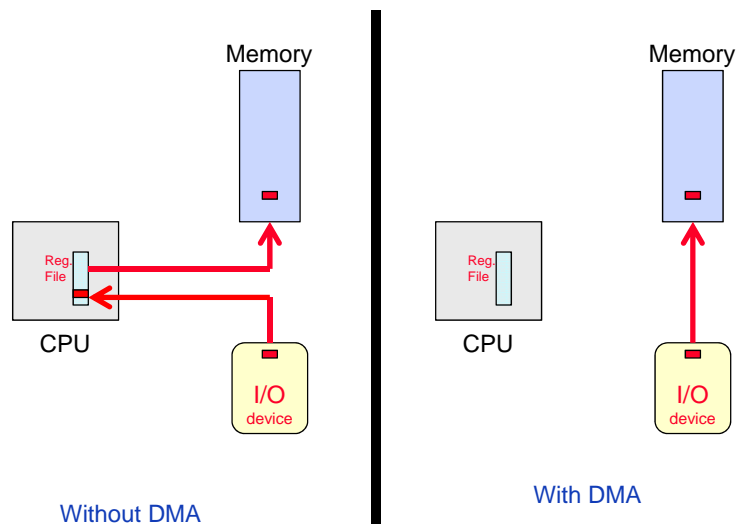
## Direct Memory Access (DMA)

- ❑ For high-bandwidth devices (like disks) polling or interrupt-driven I/O would consume a *lot* of processor cycles
- ❑ With DMA, the DMA controller has the ability to transfer large blocks of data **directly** to/from the memory without involving the processor
  1. The processor initiates the DMA transfer by supplying the I/O device address (identity), the operation to be performed, the memory address destination/source, the number of bytes to transfer
  2. The DMA controller manages the entire transfer (possibly thousand of bytes in length), arbitrating for the bus
  3. When the DMA transfer is complete (or in case of error), the DMA controller interrupts the processor to let it know that the transfer is complete
- ❑ There may be multiple DMA devices in one system
  - E.g.: systems with a single memory bus and multiple I/O buses, each I/O bus controller will often contain a DMA
  - Processor and DMA controllers contend for bus cycles and for memory
    - The processor can be delayed when the memory is busy doing a DMA transfer

Chapter 4.88

Dept. of Comp. Arch., UMA, 2018

## Direct Memory Access (DMA)



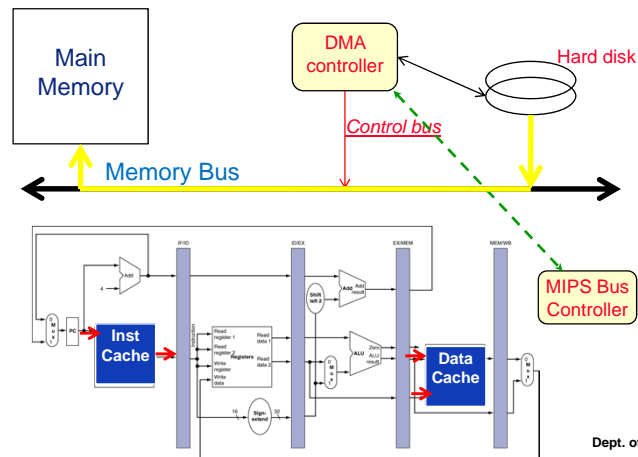
Chapter 4.89

Dept. of Comp. Arch., UMA, 2018

## Direct Memory Access (DMA)

- ❑ Processor works in parallel with the DMA controller

- Processor dealing with **Cache Buses** ( → )
- DMA controller dealing with Main **Memory Bus**



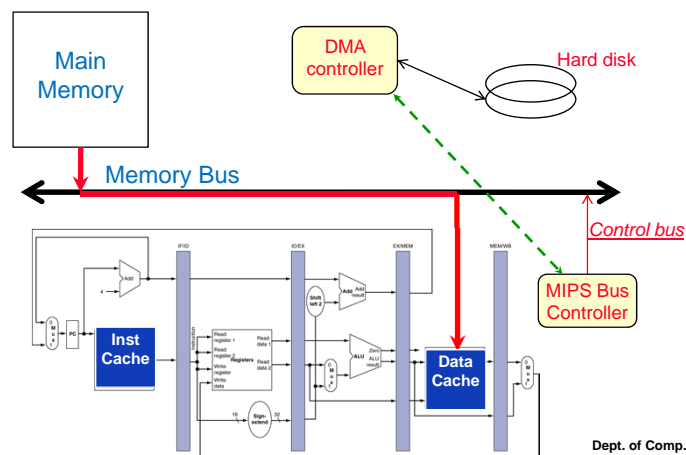
Chapter 4.90

Dept. of Comp. Arch., UMA, 2018

## Direct Memory Access (DMA)

- ❑ Example of data cache miss (or updating in a write-through)

- Processor is dealing with Main **Memory Bus**



Chapter 4.91

Dept. of Comp. Arch., UMA, 2018

## The DMA Stale Data or Coherence Problem

- ❑ In systems with caches, there can be two copies of a data item, one in the cache and one in the main memory
  - For a DMA input (from disk to memory) – the processor will be using **stale** data if that location is also in the cache
  - For a DMA output (from memory to disk) and a write-back cache – the I/O device will receive **stale** data if the data is in the cache and has not yet been written back to the memory
- ❑ The coherency problem can be solved by
  1. Routing all I/O activity through the cache – expensive and a large negative performance impact
  2. Having the OS invalidate all the entries in the cache for an I/O input or force write-backs for an I/O output (called a cache **flush**)
  3. Providing hardware to *selectively* invalidate cache entries – i.e., need a **snooping** cache controller

## DMA and Virtual Memory Considerations

- ❑ Should the DMA work with virtual addresses or physical addresses?
- ❑ If working with physical addresses
  - Must constrain all of the DMA transfers to stay within one page because if it crosses a page boundary, then it won't necessarily be contiguous in memory
  - If the transfer won't fit in a single page, it can be broken into a series of transfers (each of which fit in a page) which are handled individually and *chained* together
- ❑ If working with virtual addresses
  - The DMA controller will have to translate the virtual address to a physical address (i.e., will need a TLB structure)
- ❑ Whichever is used, the OS must cooperate by not remapping pages while a DMA transfer involving that page is in progress

## More intelligent controllers: I/O processors

- ❑ To further reduce the need to interrupt the processor the I/O controller can be made more intelligent: *I/O processors (I/O controllers or channel controllers)*
  - they execute a series of I/O operations (*I/O program* is stored in the I/O processor or in memory and fetched by the I/O processor) and interrupts the processor only when the entire program is completed.
  - The I/O program is set up by the OS: I/O operations to be done, the size and transfer address for any reads or writes
- ❑ DMA processors are essentially special-purpose processors (single-chip and nonprogrammable), while I/O processors are often implemented with general-purpose microprocessors, which run a specialized I/O program

## Interfacing I/O Devices to the Processor, Memory, and OS

- ❑ The operating system acts as the interface between the I/O hardware and the program requesting I/O
  - It provides equitable access to the shared I/O resources, protects those I/O devices/activities to which a user program doesn't have access, and schedules I/O requests to enhance system throughput
  - It handles interrupts generated by I/O devices
  - It supplies routines for low-level I/O device operations
    - OS must be able to give commands to the I/O devices
    - I/O device must be able to notify the OS about its status
    - Must be able to transfer data between the memory and the I/O device
- ❑ Software that communicates with an I/O device is called a **device driver**, and requires detailed knowledge about the I/O device hardware (ports list and its behavior)